

We will start in 5min

Getting Started with TensorFlow 2.0

TensorFlow 2.0

TensorFlow

An open source Deep Learning library

- >1,800 contributors worldwide
- Apache 2.0 license
- Released by Google in 2015

TensorFlow 2.0

- Easier to learn and use
- For beginners and experts
- Available for everyone

You can install it today

```
!pip install tensorflow
```


Amazing things (1 of 3)

Art and science in one. Deep learning as representation learning.



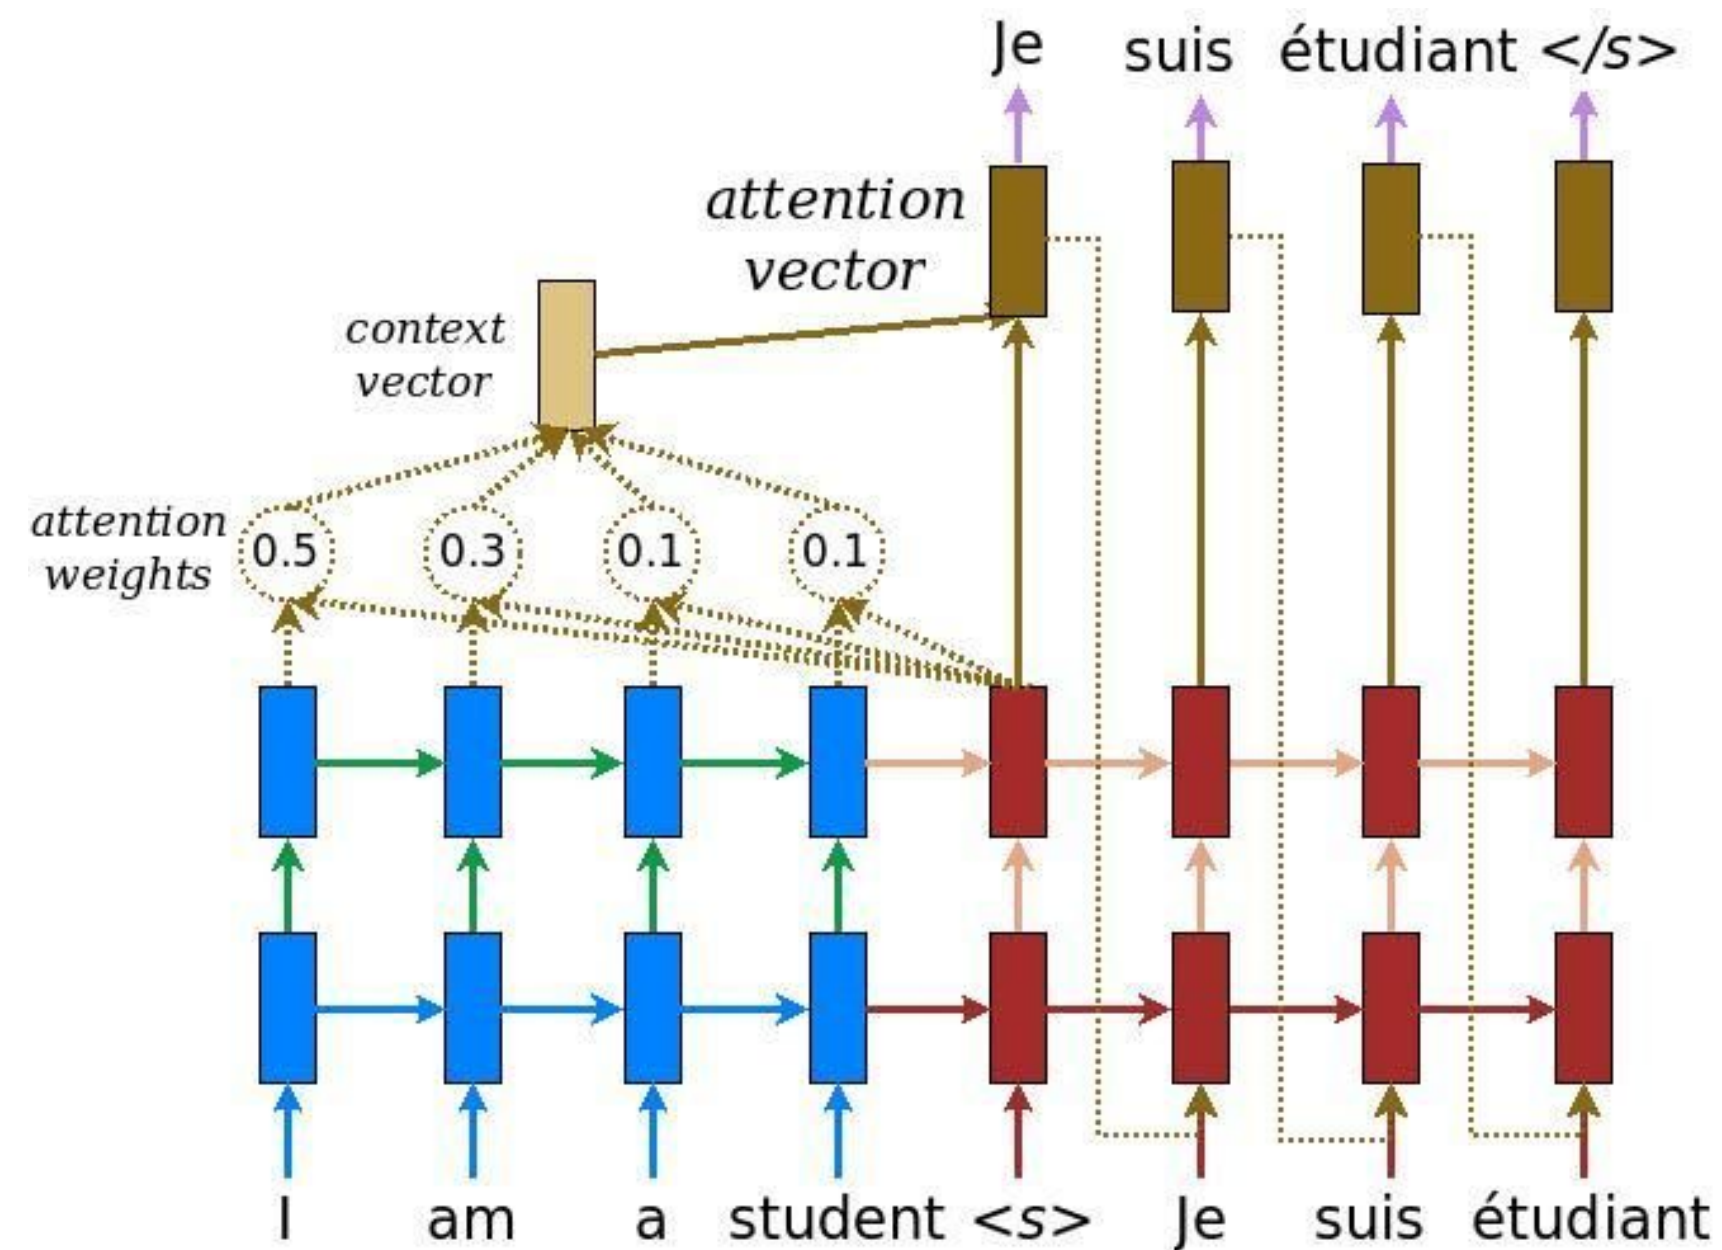
bit.ly/mini-dream

tensorflow.org/alpha/tutorials/generative/style_transfer

Amazing things (2 of 3)

Encoder / decoders. Deep learning as compression.

- Sentence -> RNN -> vector



sentences = [

("Do you want a cup of coffee?",
"¿Quieres una taza de café?"),

...

]

bit.ly/mini-nmt

tensorflow.org/alpha/tutorials/text/nmt_with_attention

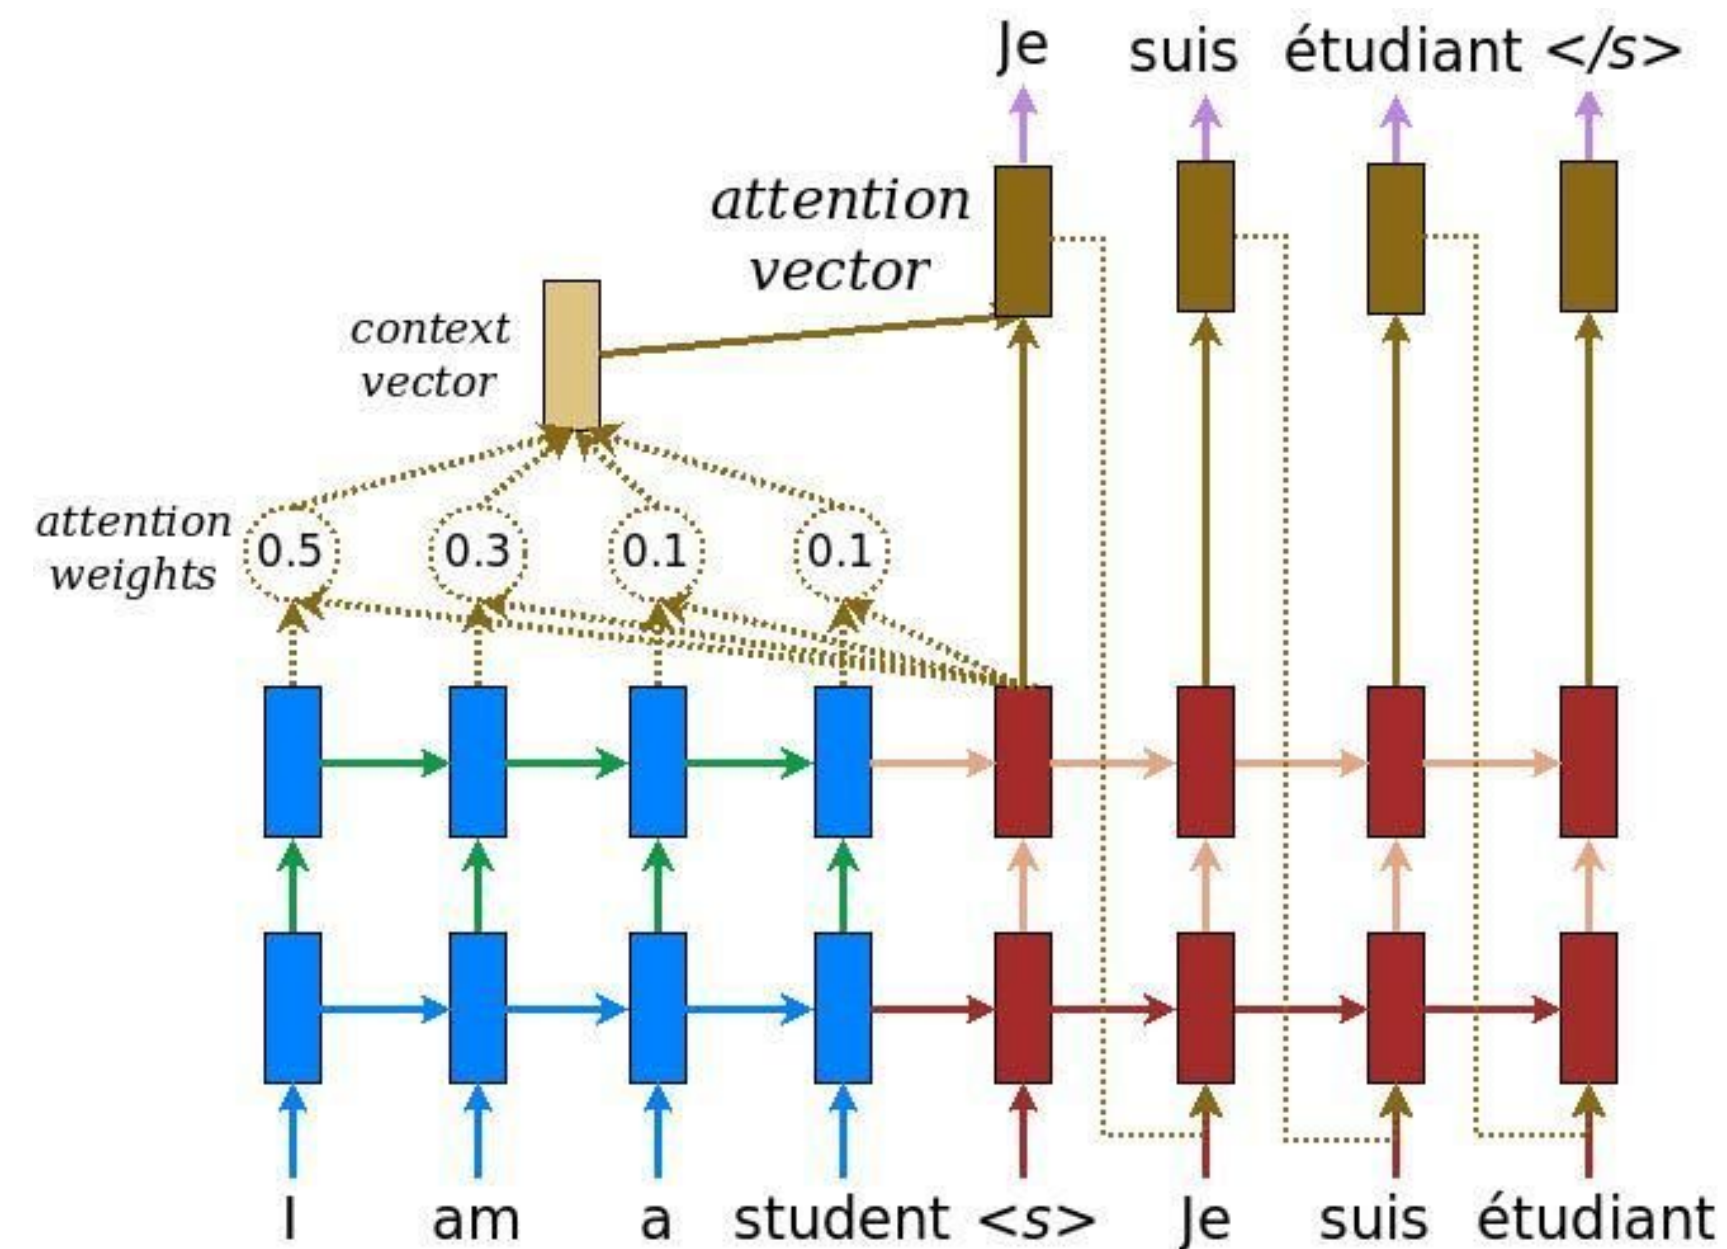
Is anyone trilingual?

Do you translate directly from **source -> target**, or from **source -> meaning -> target**?



Amazing things (2 of 3)

Opportunities at the intersection of linguistics and Deep Learning.



sentences = [

("Do you want a cup of coffee?",
"¿Quieres una taza de café?"),

...

("Voulez-vous une tasse de café?",
"¿Quieres una taza de café?")



...

]

bit.ly/mini-nmt

tensorflow.org/alpha/tutorials/text/nmt_with_attention

TensorFlow is great for learning (with 2.0, it's a perfect time to start)



bit.ly/tf-linear

[tensorflow.org/alpha/tutorials/quickstart/beginner](https://www.tensorflow.org/alpha/tutorials/quickstart/beginner)


```
x = tf.placeholder(tf.float32, [None, 200])
```

```
W = tf.Variable(tf.zeros([200, 10]))
```

```
b = tf.Variable(tf.zeros([10]))
```

```
y = tf.nn.softmax(tf.matmul(x, W) + b)
```

```
...
```

```
with tf.Session() as sess:
```

```
    sess.run(tf.initialize_all_variables())
```

```
    tf.train.start_queue_runners(sess)
```

```
example_batch = tf.train.batch([x], batch_size=10, num_threads=4, capacity=10)
```

```
max_steps = 1000
```

```
for step in range(max_steps):
```

```
    x_in = sess.run(example_batch)
```

```
    sess.run(train_step, feed_dict={x: train_data, y_: train_labels})
```

```
    if (step % 100) == 0:
```

```
        print(step, sess.run(accuracy, feed_dict={x: test_data, y_: test_labels}))
```

We've learned a lot
since 1.0



TensorFlow 2.0

Usability

- `tf.keras` as the recommended high-level API.
- Eager execution by default.

```
>>> tf.add(2, 3)
<tf.Tensor: id=2, shape=(), dtype=int32, numpy=5>
```



TensorFlow 2.0

Clarity

-
- Remove duplicate functionality
 - Consistent, intuitive syntax across APIs
 - Compatibility throughout the TensorFlow ecosystem




TensorFlow 2.0

Flexibility

-
- Full lower-level API.
 - Internal ops accessible in `tf.raw_ops`
 - Inheritable interfaces for variables, checkpoints, layers.

All with one API



With styles for
beginners and experts

For beginners

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

TF 1.x

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```


TF 2.0

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

For experts

```
class MyModel(tf.keras.Model):  
    def __init__(self, num_classes=10):  
        super(MyModel, self).__init__(name='my_model')  
        self.dense_1 = layers.Dense(32, activation='relu')  
        self.dense_2 = layers.Dense(num_classes, activation='sigmoid')  
  
    def call(self, inputs):  
        # Define your forward pass here,  
        x = self.dense_1(inputs)  
        return self.dense_2(x)
```

Technical differences

Symbolic (Sequential)

- Your model is a graph of layers
- Any model you compile will run
- TensorFlow **helps you debug** by catching errors at **compile time**

Technical differences

Symbolic (Sequential)

- Your model is a graph of layers
- Any model you compile will run
- TensorFlow **helps you debug** by catching errors at **compile time**

Imperative (Subclassing)

- Your model is Python bytecode
- Complete flexibility and control
- Harder to debug / **harder to maintain**



Training loops

Use a built-in training loop...

```
model.fit(x_train, y_train, epochs=5)
```

Or define your own

```
model = MyModel()

with tf.GradientTape() as tape:
    logits = model(images)
    loss_value = loss(logits, labels)

grads = tape.gradient(loss_value, model.trainable_variables)
optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

Keras vs tf.keras?

keras.io

- Reference implementation
- *import keras*

tf.keras

- TensorFlow's implementation (a superset, built-in to TF)
- *from tensorflow import keras*

Going big: tf.distribute.Strategy

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, input_shape=[10]),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')])

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

Going big: Multi-GPU

```
strategy = tf.distribute.MirroredStrategy()

with strategy.scope():
    model = tf.keras.models.Sequential([
        tf.keras.layers.Dense(64, input_shape=[10]),
        tf.keras.layers.Dense(64, activation='relu'),
        tf.keras.layers.Dense(10, activation='softmax')])

    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```




The docs are runnable

Pix2Pix

Speaking of amazing things

 Run in Google Colab

 View source on GitHub

This notebook demonstrates image to image translation using conditional GAN's, as described in [Image-to-Image Translation with Conditional Adversarial Networks](#). Using this technique we can colorize black and white photos, convert google maps to google earth, etc. Here, we convert building facades to real buildings.



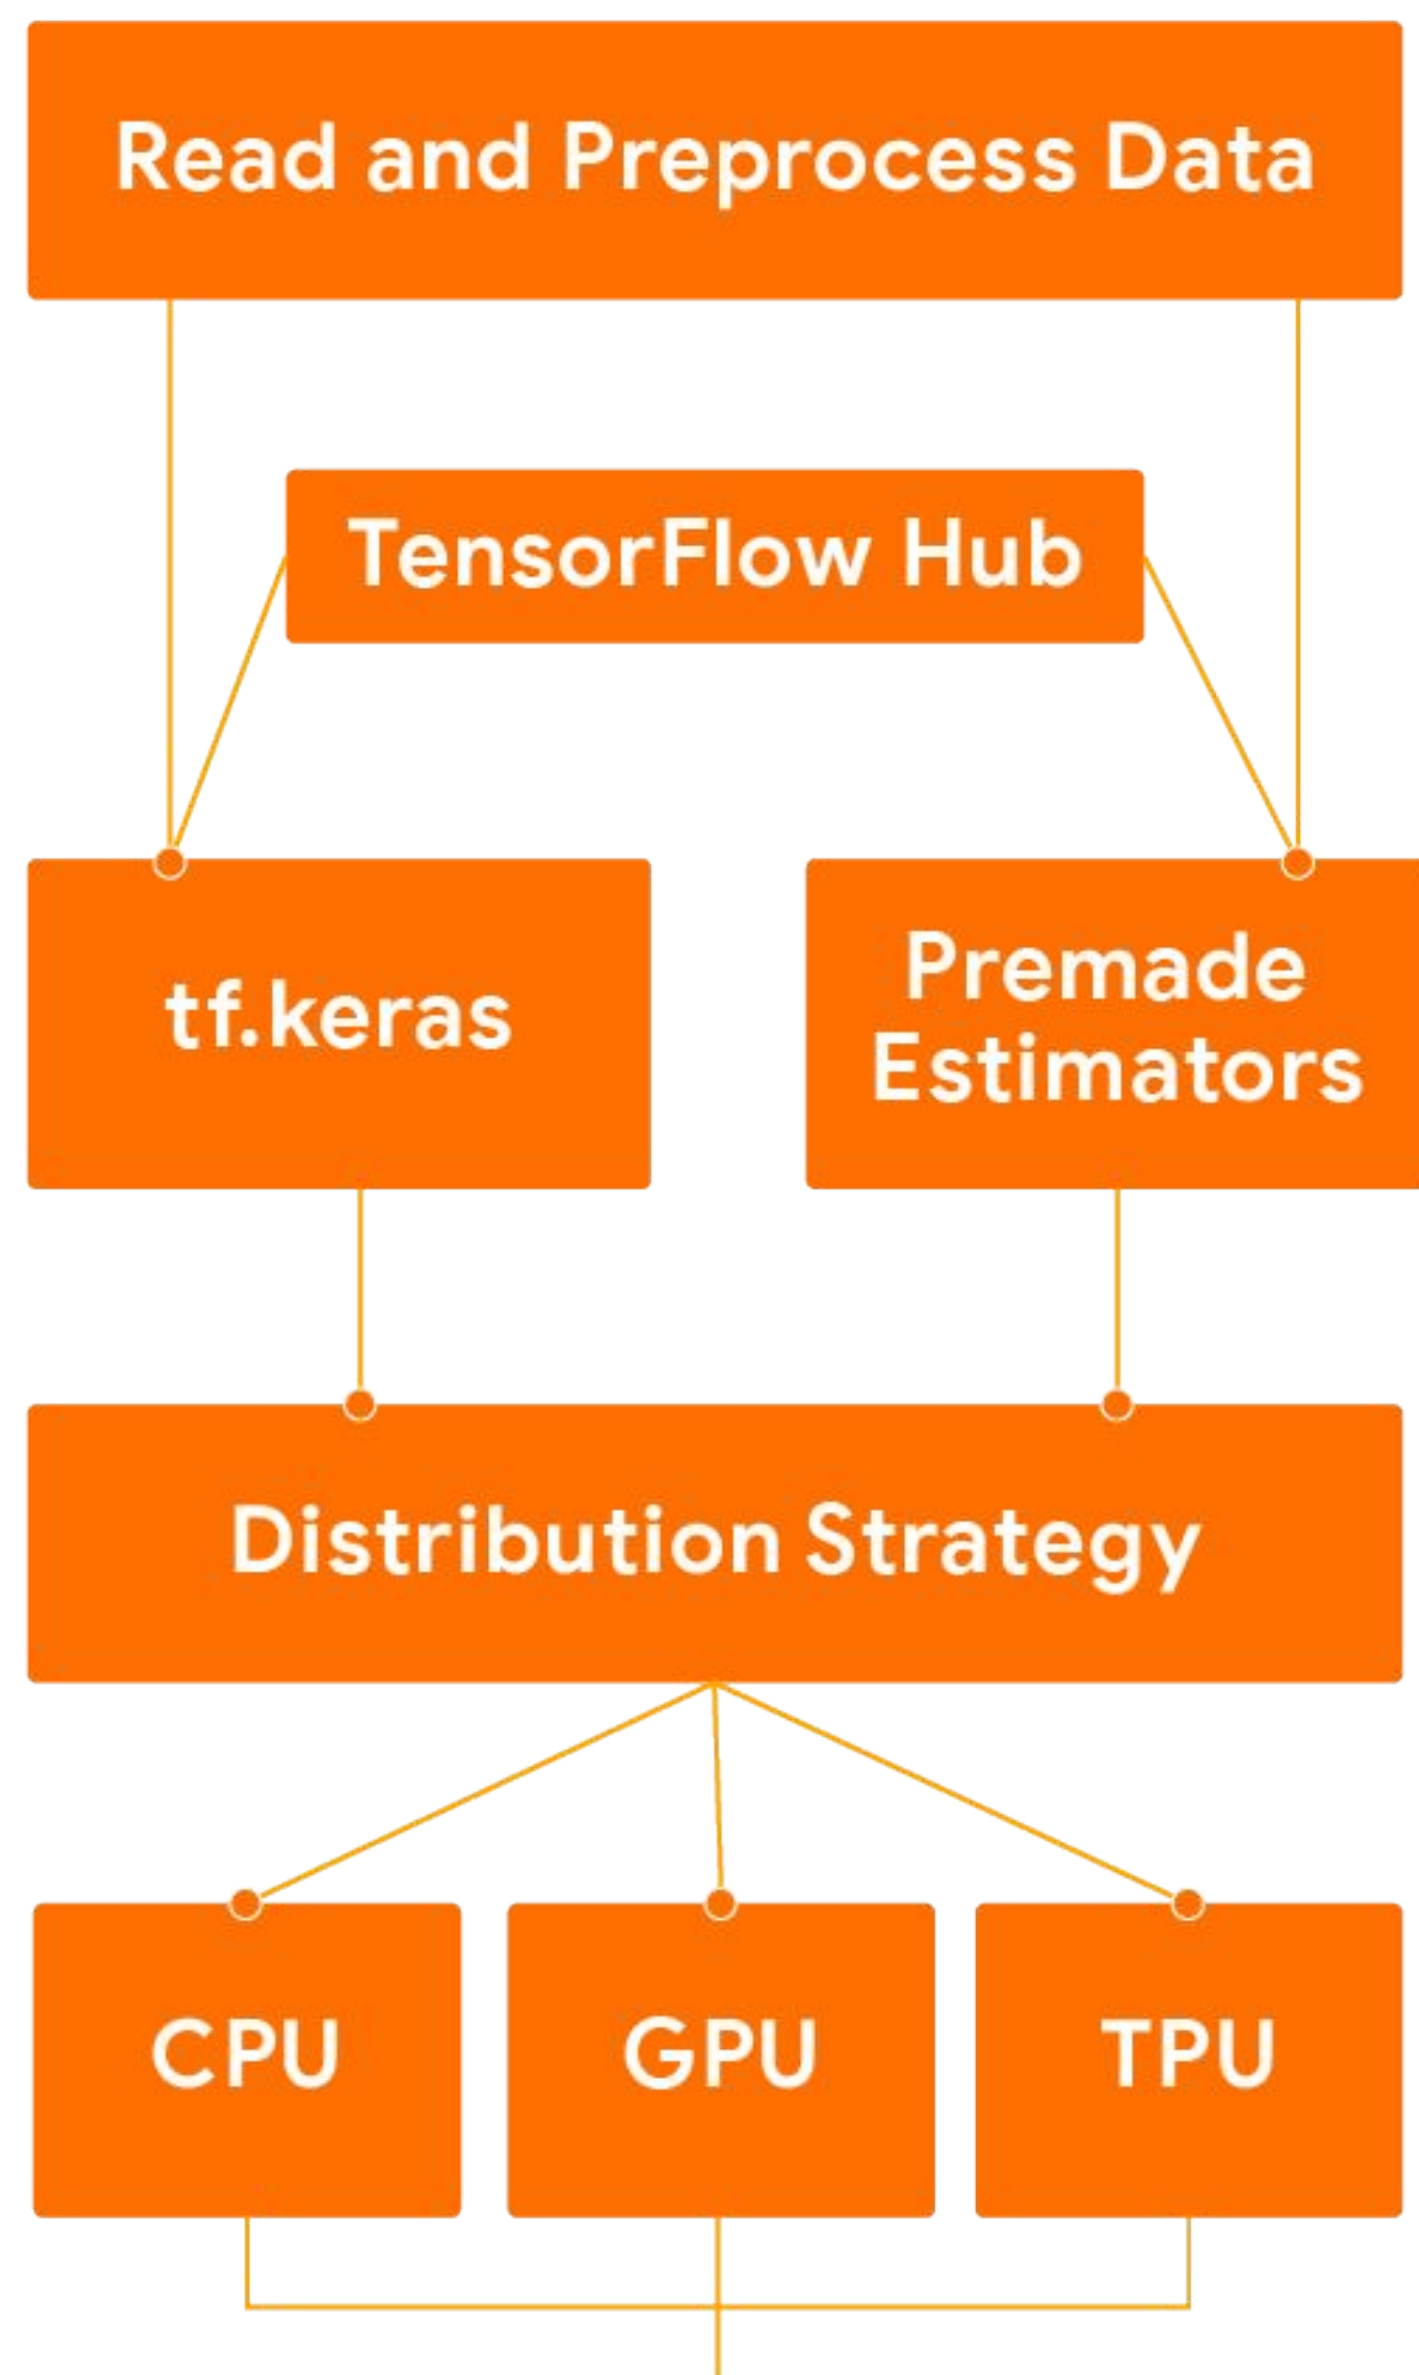
tensorflow.org/alpha

tensorflow.org/alpha/tutorials/generative/pix2pix

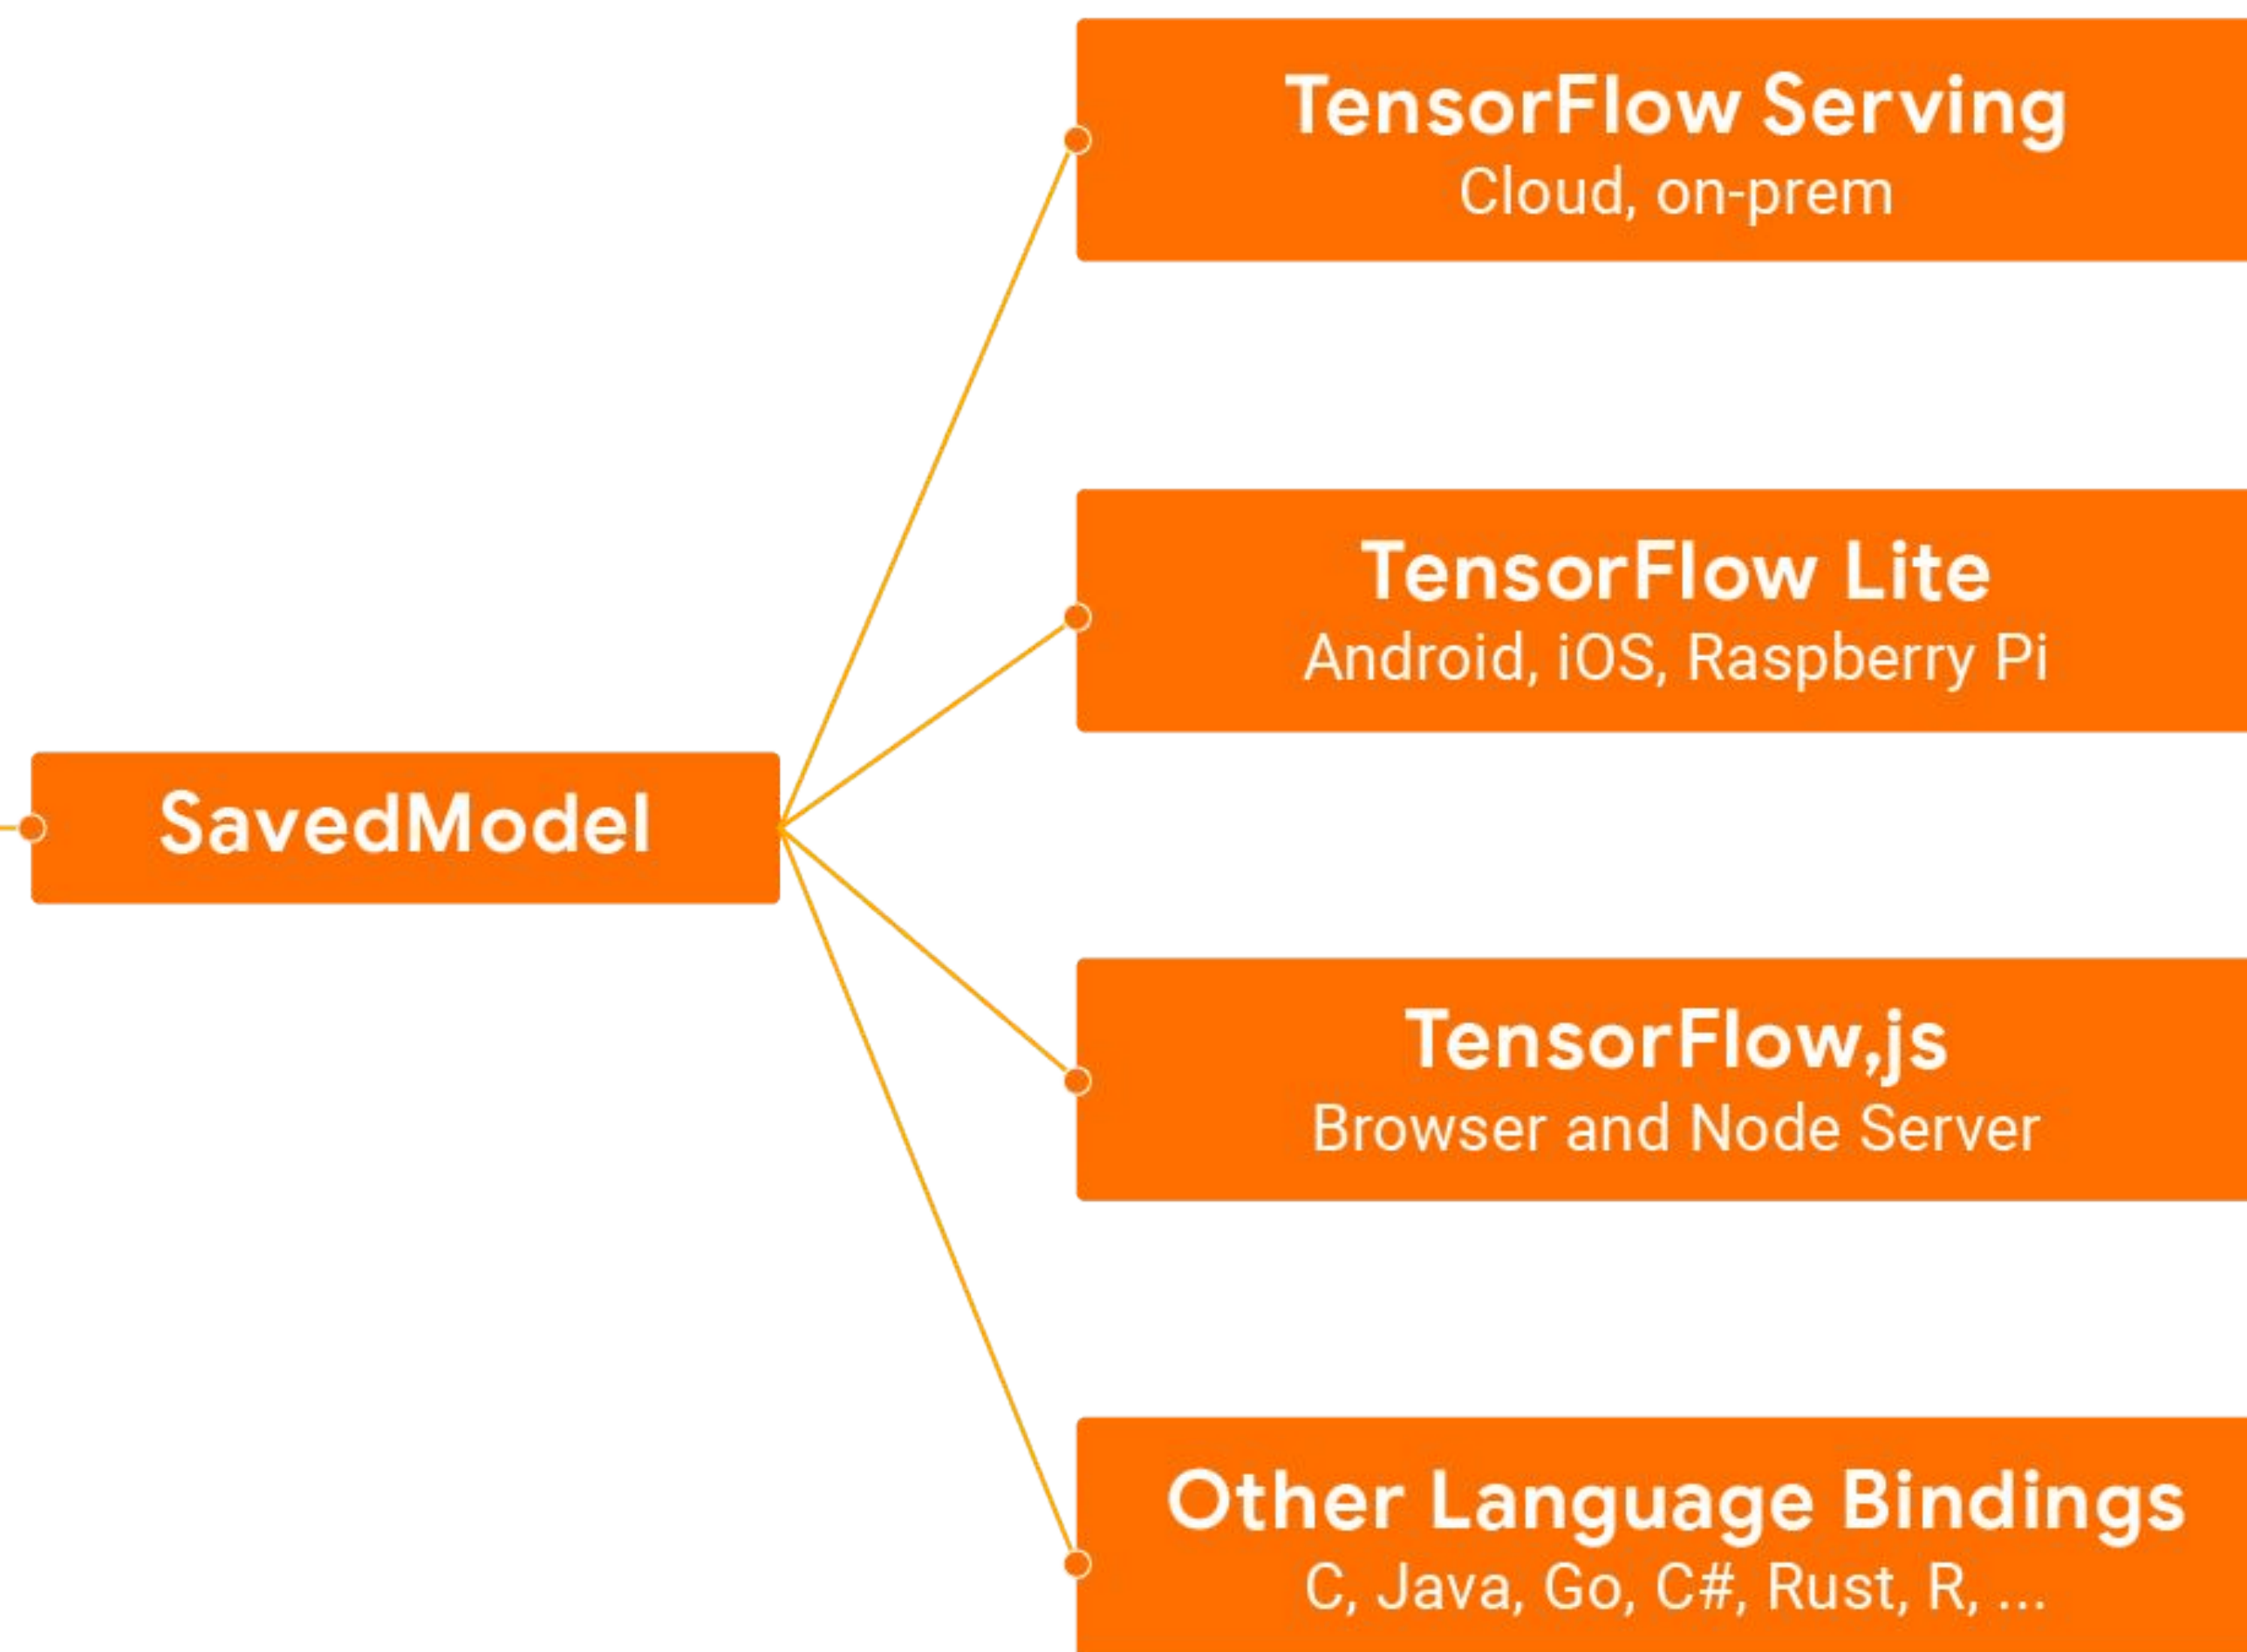
Compatibility with the TensorFlow Ecosystem



TRAINING



DEPLOYMENT





Deploy anywhere

Servers



TensorFlow
Extended

Edge devices



TensorFlow
Lite

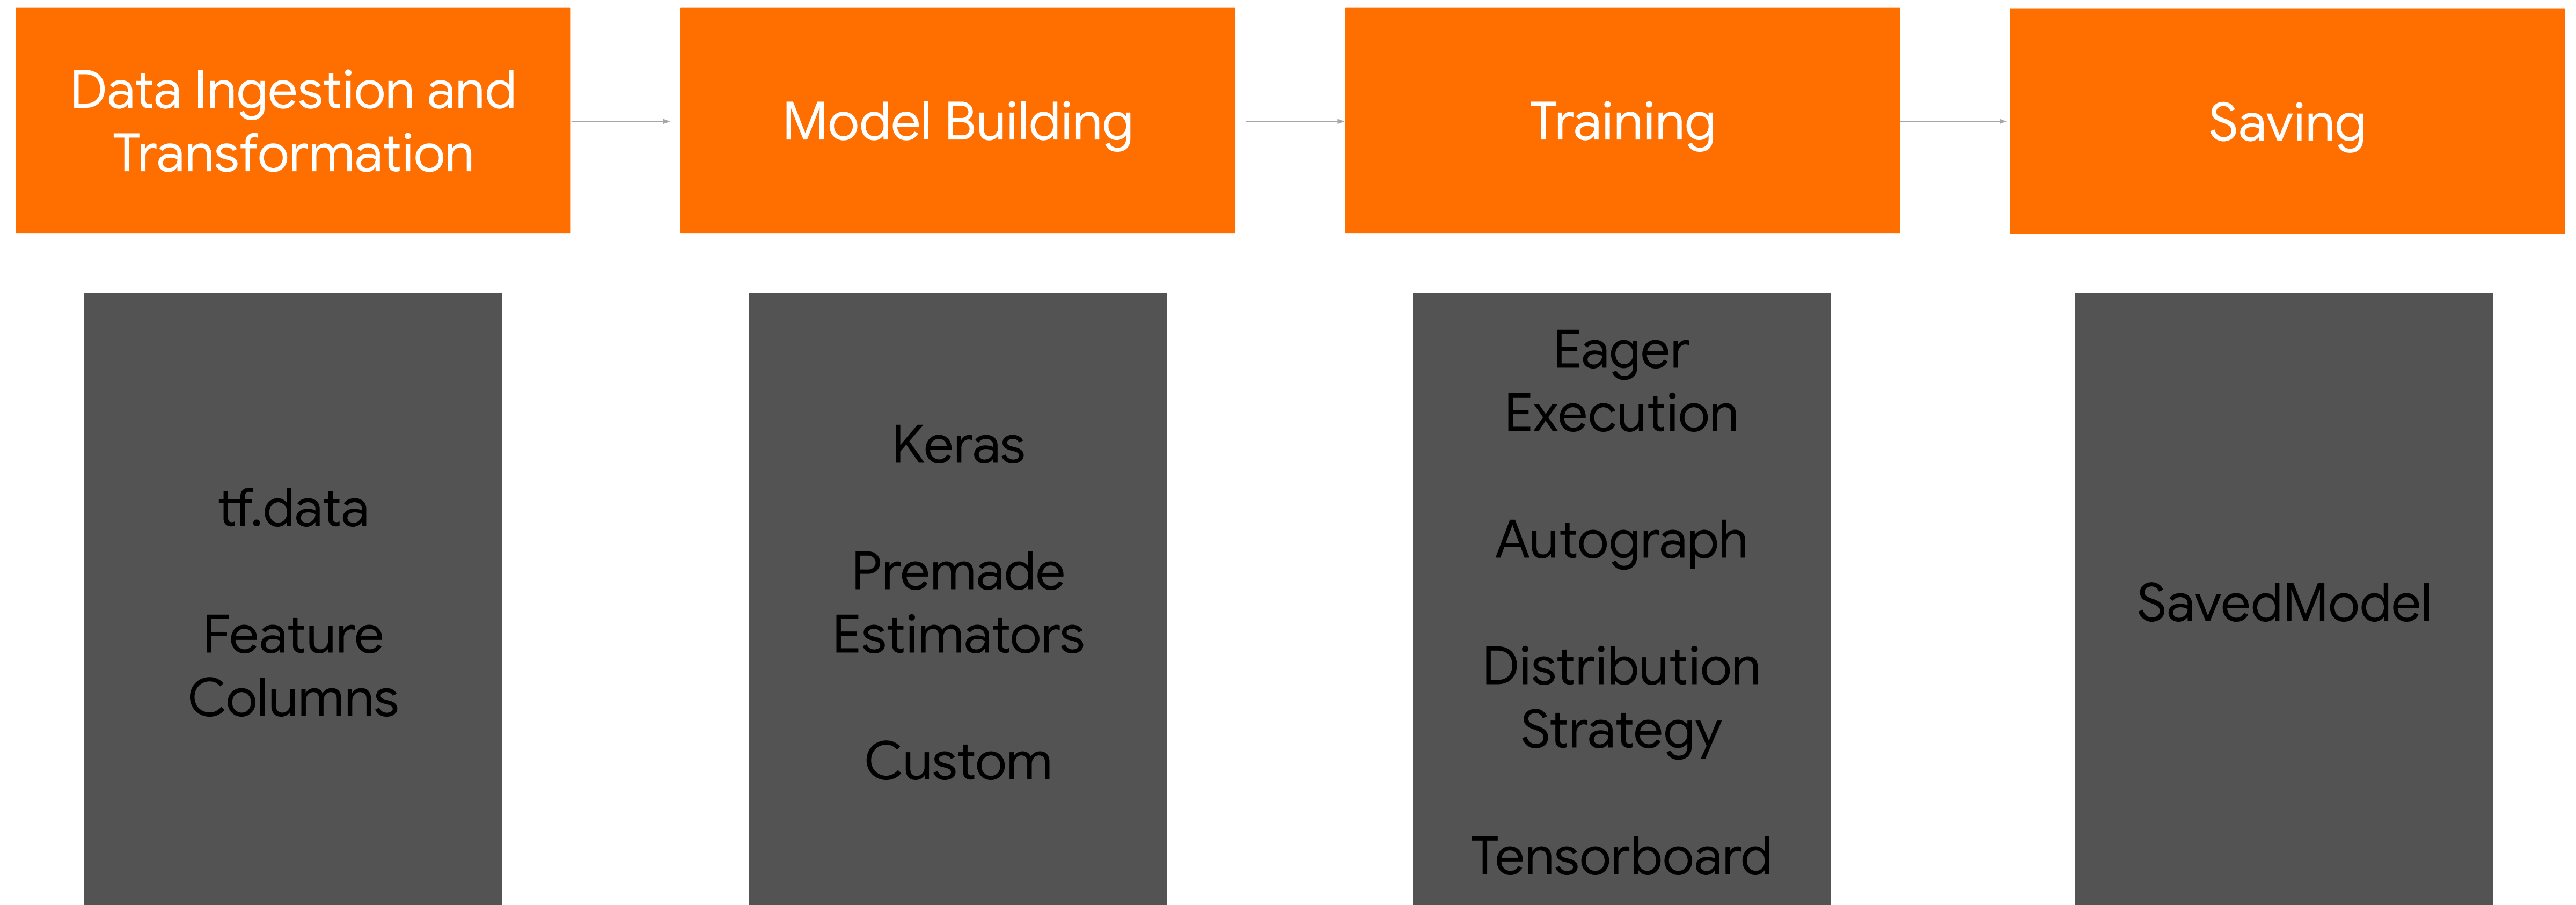
JavaScript



TensorFlow
.JS



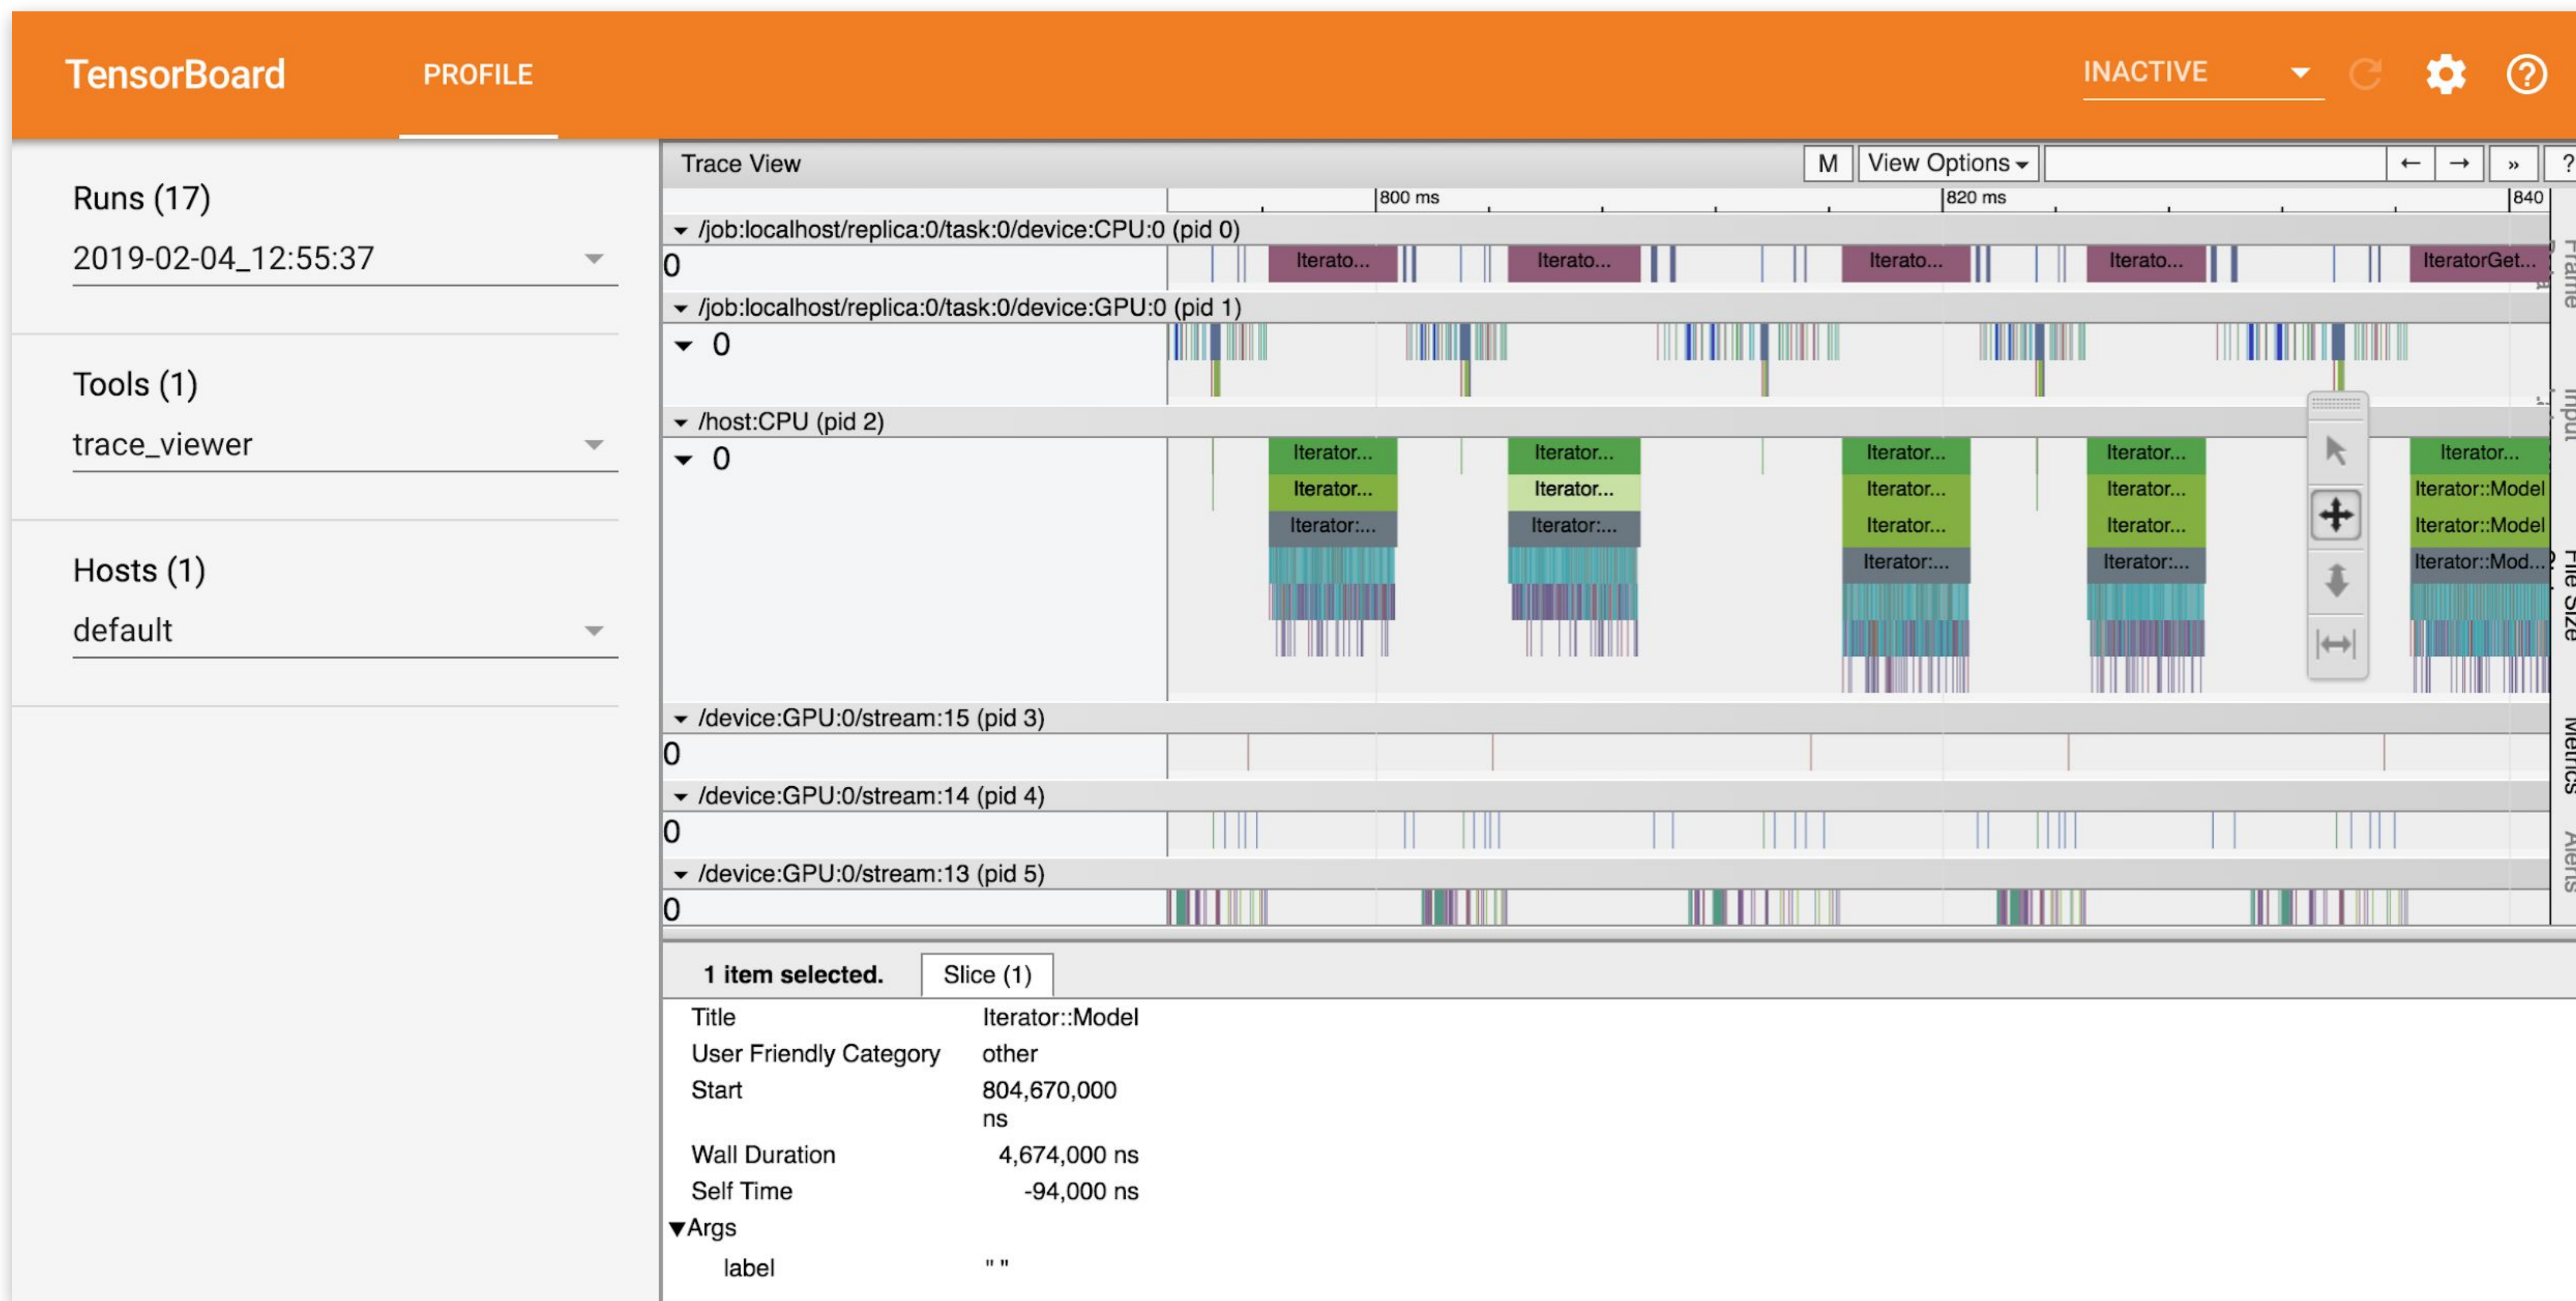
Training Workflow







Built-in performance profiling





TensorFlow

Install

Learn ▾

API ▾

Resources ▾

Community

Why TensorFlow ▾

Search

GitHub

API r2.0

Python

JavaScript

C++

Java

More...

Python r2.0

tf

Overview

AggregationMethod

argsort

batch_to_space

bitcast

boolean_mask

broadcast_dynamic_shape

broadcast_static_shape

broadcast_to

case

clip_by_global_norm

clip_by_norm

clip_by_value

combined_non_max_suppression

concat

cond

constant

constant_initializer

control_dependencies

convert_to_tensor

CriticalSection

custom_gradient

device

div_no_nan

dynamic_partition

Announcing the Tensorflow Dev Summit 2019

Learn More

TensorFlow > API r2.0 > Python

☆☆☆☆☆

Module: tf

Defined in `__init__.py`.

Bring in all of the public TensorFlow interface into this module.

Modules

`audio` module: Public API for tf.audio namespace.

`autograph` module: Conversion of plain Python into TensorFlow graph code.

`bitwise` module: Operations for manipulating the binary representations of integers.

`compat` module: Functions for Python 2 vs. 3 compatibility.

`config` module: Public API for tf.config namespace.

`data` module: `tf.data.Dataset` API for input pipelines.

`debugging` module: Public API for tf.debugging namespace.

Contents

Modules

Classes

Functions

Other Members

`__all__`

`__compiler_version__`

`__cxx11_abi_flag__`

`__git_version__`

`__monolithic_build__`

`__version__`

`bfloat16`

`bool`

`complex128`

`complex64`

`double`

`float16`

`float32`

`float64`

`half`

`int16`

`int32`

`int64`

`int8`

`newaxis`

`plugin_dir`

`qint16`



Developers need great performance

Since last year...

1.8x training speedup on NVIDIA Tesla V100

1.6x training speedup on Google Cloud TPU v2

3.3x inference speedup on Intel Skylake



Focused on speed

Incredible inference performance

CPU
124 ms

CPU on
MobileNet V1

CPU 1.9x
64ms

CPU w/
Quantization

GPU 7.7x
16 ms

Flow
OpenGL 16

Edge TPU 62x
2 ms

Quantized
Fixed-point

Pixel 2 - Single Threaded CPU



TF Probability

TF Agents

Tensor2Tensor

TF Ranking

TF Text

TF Federated

TF Privacy

...



How do I upgrade?



Making upgrading easy:

- Escape to backwards compatibility module:
`tf.compat.v1` (does not include `tf.contrib`)
- Migration guides and best practices
- Conversion script: `tf_upgrade_v2`



Copyright 2018 The TensorFlow Authors.

Licensed under the Apache License, Version 2.0 (the "License");

Text generation using a RNN with eager execution

Setup

Import TensorFlow and other libraries

Download the Shakespeare dataset

Read the data

Process the text

Vectorize the text

The prediction task

Create training examples and targets

Create training batches

Build The Model

Try the model

Train the model

Attach an optimizer, and a loss function

Configure checkpoints

▼ Copyright 2018 The TensorFlow Authors.

[] Licensed under the Apache License, Version 2.0 (the "License");

▼ Text generation using a RNN with eager execution

[View on TensorFlow.org](#)[Run in Google Colab](#)[View source on GitHub](#)

[] !pip install tf-nightly

This tutorial demonstrates how to generate text using a character-based RNN. We will work with a dataset of Shakespeare's writing from Andrej Karpathy's [The Unreasonable Effectiveness of Recurrent Neural Networks](#). Given a sequence of characters from this data ("Shakespear"), train a model to predict the next character in the sequence ("e"). Longer sequences of text can be generated by calling the model repeatedly.

Note: Enable GPU acceleration to execute this notebook faster. In Colab: *Runtime > Change runtime type > Hardware accelerator > GPU*. If running locally make sure TensorFlow version ≥ 1.11 .

This tutorial includes runnable code implemented using [tf.keras](#) and [eager execution](#). The following is sample output when the model in this tutorial trained for 30 epochs, and started with the string "Q":

QUEENE:

I had thought thou hadst a Roman; for the oracle,
Thus by All bids the man against the word,
Which are so weak of care, by old care done;
Your children were in your holy love,
And the precipitation through the bleeding throne.

BISHOP OF ELY:

Marry, and will, my lord, to weep in such a one were prettiest;
Yet now I was adopted heir
Of the world's lamentable day,



Table of contents Code snippets Files X

↑ UPLOAD ↻ REFRESH



..

sample_data

text_generation.py

```
[2] !pip install tf-nightly-2.0-preview
```

```
[3] import tensorflow as tf
     tf.__version__
```

```
↳ '2.0.0-dev20190218'
```

```
⌂ !tf_upgrade_v2 --infile text_generation.py --outfile text_generation_upgraded.py
```

```
▶ !cat report.txt
```

```
[ ]
```



```
[4] !tf_upgrade_v2 --infile text_generation.py --outfile text_generation_upgraded.py
```

```
[>] INFO line 4:0: Renamed 'tf.enable_eager_execution' to 'tf.compat.v1.enable_eager_execution'
INFO line 240:16: Renamed 'tf.train.AdamOptimizer' to 'tf.compat.v1.train.AdamOptimizer'
INFO line 332:21: Added keywords to args of function 'tf.multinomial'
INFO line 332:21: Renamed 'tf.multinomial' to 'tf.random.categorical'
INFO line 375:12: Renamed 'tf.train.AdamOptimizer' to 'tf.compat.v1.train.AdamOptimizer'
INFO line 392:21: tf.losses.sparse_softmax_cross_entropy requires manual check. tf.losses.sparse_softmax_cross_entropy
INFO line 392:21: Renamed 'tf.losses.sparse_softmax_cross_entropy' to 'tf.compat.v1.losses.sparse_softmax_cross_entropy'
TensorFlow 2.0 Upgrade Script
-----
Converted 1 files
Detected 0 issues that require attention
-----
```

Make sure to read the detailed log 'report.txt'



```
[4] !tf_upgrade_v2 --infile text_generation.py --outfile text_generation_upgraded.py
```

```
[>] INFO line 4:0: Renamed 'tf.enable_eager_execution' to 'tf.compat.v1.enable_eager_execution'
INFO line 240:16: Renamed 'tf.train.AdamOptimizer' to 'tf.compat.v1.train.AdamOptimizer'
INFO line 332:21: Added keywords to args of function 'tf.multinomial'
INFO line 332:21: Renamed 'tf.multinomial' to 'tf.random.categorical'
INFO line 375:12: Renamed 'tf.train.AdamOptimizer' to 'tf.compat.v1.train.AdamOptimizer'
INFO line 392:21: tf.losses.sparse_softmax_cross_entropy requires manual check. tf.losses.sparse_softmax_cross_entropy
INFO line 392:21: Renamed 'tf.losses.sparse_softmax_cross_entropy' to 'tf.compat.v1.losses.sparse_softmax_cross_entropy'
TensorFlow 2.0 Upgrade Script
-----
Converted 1 files
Detected 0 issues that require attention
-----

Make sure to read the detailed log 'report.txt'
```




```
[4] !tf_upgrade_v2 --infile text_generation.py --outfile text_generation_upgraded.py
```

```
[>] INFO line 4:0: Renamed 'tf.enable_eager_execution' to 'tf.compat.v1.enable_eager_execution'
INFO line 240:16: Renamed 'tf.train.AdamOptimizer' to 'tf.compat.v1.train.AdamOptimizer'
INFO line 332:21: Added keywords to args of function 'tf.multinomial'
INFO line 332:21: Renamed 'tf.multinomial' to 'tf.random.categorical'
INFO line 375:12: Renamed 'tf.train.AdamOptimizer' to 'tf.compat.v1.train.AdamOptimizer'
INFO line 392:21: tf.losses.sparse_softmax_cross_entropy requires manual check. tf.losses.sparse_softmax_cross_entropy
INFO line 392:21: Renamed 'tf.losses.sparse_softmax_cross_entropy' to 'tf.compat.v1.losses.sparse_softmax_cross_entropy'
TensorFlow 2.0 Upgrade Script
-----
Converted 1 files
Detected 0 issues that require attention
-----
```

Make sure to read the detailed log 'report.txt'

Table of contents Code snippets **Files** ✕

↑ UPLOAD ↻ REFRESH



..

sample_data

report.txt

text_generation.py

text_generation_upgraded.py

[4]



```
=====
Processing file 'text_generation.py'
outputting to 'text_generation_upgraded.py'
=====
```

```
4:0: INFO: Renamed 'tf.enable_eager_execution' to 'tf.compat.v1.enable_eager_execution'
```

```
240:16: INFO: Renamed 'tf.train.AdamOptimizer' to 'tf.compat.v1.train.AdamOptimizer'
```

```
332:21: INFO: Added keywords to args of function 'tf.multinomial'
```

```
332:21: INFO: Renamed 'tf.multinomial' to 'tf.random.categorical'
```

```
375:12: INFO: Renamed 'tf.train.AdamOptimizer' to 'tf.compat.v1.train.AdamOptimizer'
```

```
392:21: INFO: tf.losses.sparse_softmax_cross_entropy requires manual check. tf.losses have been replaced with object
```

```
392:21: INFO: Renamed 'tf.losses.sparse_softmax_cross_entropy' to 'tf.compat.v1.losses.sparse_softmax_cross_entropy'
```



|





Search or jump to... /

Pull requests

Issues

Marketplace

Explore

lc0 / deeplearning-playground

Unwatch

1

Star

0

Fork

0

<> Code

Issues 0

Pull requests 0

Projects 0

Wiki

Insights

Settings

Experiments and ideas in notebooks

Edit

playground

tensorflow

Manage topics

11 commits

1 branch

0 releases

1 contributor

Branch: master

New pull request

Create new file

Upload files

Find file

Clone or download

lc0 Update minor fixes

Latest commit 0006856 18 days ago

TF2	Update minor fixes	18 days ago
charts	Updated scale for graphs	3 months ago
.gitignore	Add an example for Fashion MNIST on TPU with keras	5 months ago
Bucharest_AI_Shap_Values.ipynb	Add SHAP values for fashion MNIST	4 months ago
Serving_REST_simple.ipynb	Add example with TF Serving	3 months ago
TPU_Keras_Fashion_MNIST.ipynb	Add an example for Fashion MNIST on TPU with keras	5 months ago
TensorFlow_Ranking.ipynb	Add example to Setup TF Ranking in Colab	3 months ago
The_simplest_fashion_mnist+_Confusio...	Test save from Google Colaboratory	3 months ago
adversarial.ipynb	Robust rotation	2 years ago

Help people interested in this repository understand your project by adding a README.

Add a README

Elements

Console

Sources

Network

Performance

Memory

Application

Security

Audits

AdBlock

op

Filter

Default levels

http://tf2up.ml



Timeline

Now

Spring

2.0 alpha
cut from head, updated frequently

2.0 RC
(branch created)

2.0
final

- implement remaining features
- convert libraries, Google
- lots of testing, optimization

- release & integration testing



Progress

Search or jump to... / Pull requests Issues Marketplace Explore

tensorflow

Repositories 87 People 476 Teams 54 Projects 2 Settings

TensorFlow 2.0
Updated 4 days ago

Filter cards

5 To do + ...

Feature: Unify argument names (~70). ...
tensorflow#25357 opened by dynamicwebpaige
2.0 type:feature

Feature: Support for FP16 training in Keras. ...
tensorflow#25359 opened by dynamicwebpaige
2.0 type:feature

Feature: Update the environment capture script for system information. ...
tensorflow#25461 opened by dynamicwebpaige
2.0 stat:contributions ... type:feature
type:support

39 In progress + ...

[TF2.0] Custom variables ...
tensorflow#24865 opened by awav
2.0

Tensorflow 2.0 Preview - tf.function and tensorflow_dataset incompatibility ...
tensorflow#25414 opened by EmanueleGhelfi
2.0 comp:data

TFLite Converter not able to convert tf.keras model in TensorFlow 2.0 ...
tensorflow#25575 opened by margaretmz
2.0 comp:lite

tf.rnn wrappers are incompatible with tf.keras.layers cells ...
tensorflow#25611 opened by dynamicwebpaige
2.0 type:bug

7 Done + ...

TF 2.0: tf.estimator import issue. ...
tensorflow#25723 opened by dynamicwebpaige
2.0 comp:apis

TF 2.0: Create conversion tool from TensorFlow 1.x. ...
tensorflow#25371 opened by dynamicwebpaige
2.0 type:support

TF 2.0: Create nightly CI builds for the TF 2.0 target. ...
tensorflow#25370 opened by dynamicwebpaige
2.0 type:build/install

TF 2.0: API symbol renames design review ...
tensorflow#25369 opened by dynamicwebpaige
2.0 type:design

Under the hood

What exactly *is* TensorFlow?
What problems does it solve?

About how much slower is Python than C?

- Multiplying matrices: +/- 100X
- 6 seconds vs. 10 minutes
- Running vs. flying (6 MPH and 600 MPH)

Python is *incredibly popular* for scientific computing

- Why? NumPy!
- C performance, Python ease of use

TensorFlow is a similar idea

- A C++ backend with a Python frontend

You can use TF 2.0 like NumPy

```
import tensorflow as tf # Assuming TF 2.0 is installed

a = tf.constant([[1, 2],[3, 4]])
b = tf.matmul(a, a)

print(b)
# tf.Tensor( [[ 7 10] [15 22]], shape=(2, 2), dtype=int32)

print(type(b.numpy()))
# <class 'numpy.ndarray'>
```


Let's make this faster

```
lstm_cell = tf.keras.layers.LSTMCell(10)
```

```
def fn(input, state):  
    return lstm_cell(input, state)
```

```
input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2  
lstm_cell(input, state); fn(input, state) # warm up
```

```
# benchmark
```

```
timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
```

Let's make this faster

```
lstm_cell = tf.keras.layers.LSTMCell(10)
```

```
@tf.function
```

```
def fn(input, state):
```

```
    return lstm_cell(input, state)
```

```
input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2
```

```
lstm_cell(input, state); fn(input, state) # warm up
```

```
# benchmark
```

```
timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
```

```
timeit.timeit(lambda: fn(input, state), number=10) # 0.004
```

AutoGraph makes this possible

```
@tf.function
def f(x):
    while tf.reduce_sum(x) > 1:
        x = tf.tanh(x)
    return x

# you never need to run this (unless curious)
print(tf.autograph.to_code(f))
```

Generated code

```
def tf__f(x):
    def loop_test(x_1):
        with ag__.function_scope('loop_test'):
            return ag__.gt(tf.reduce_sum(x_1), 1)
    def loop_body(x_1):
        with ag__.function_scope('loop_body'):
            with ag__.utils.control_dependency_on_returns(tf.print(x_1)):
                tf_1, x = ag__.utils.alias_tensors(tf, x_1)
                x = tf_1.tanh(x)
            return x,
    x = ag__.while_stmt(loop_test, loop_body, (x,), (tf,))
    return x
```


Learning more

Tutorials

tensorflow.org/beta

Tip: Ways to know you're using a 1.x tutorial

- `tf.enable_eager_execution()`
- `session.run`
- `tf.placeholder`
- `feed_dict`

Books

Two of our favorites

- [Hands-on Machine Learning with Scikit-Learn, Keras, and TensorFlow 2.0](#), by Aurélien Géron
- [Deep Learning with Python](#), by François Chollet

Thank you!



Abderrahman JAIZE

twitter.com/jaizemsd